

# Changepoint detection

*Inference and applications*

**Etienne Caprioli**

*etienne.caprioli23@imperial.ac.uk*

*Imperial College London*

June 11, 2024



# Outline of this presentation

## *Inference*

- Problem setting
  - Statistical test
  - Online/Offline inference
- What are we interested in ?
  - Single changepoint
  - Multiple changepoints
- Bayesian approach

## *Applications*

- Application in market risk monitoring
- Other areas of applications
- Useful libraries, packages and repositories

# Inference

# Problem setting

## *Piecewise homogeneous sequences*

Let  $(X_i)_i$  be a sequence of random variables. We model this signal with the equation

$$X_t = \sum_{j=0}^q X_t^{(j)} \mathbf{1}_{\tau_j \leq t \leq \tau_{j+1}}$$

Where each  $(X_t^{(j)})_j$  can be modeled individually. We are interested in inferring the *changepoints*  $(\tau_j)_j$  and the  *$j$  joint distributions* of  $(X_t^{(j)})_j$ .


### Note

The collections  $(X_t^{(j)})_j$  can be modeled in very different ways: IID sequences, (non-) stationary time series, IID + trend/seasonality...


# Problem setting

## *Example of datasets*

Change in mean



Change in variance



# Problem setting

## *Statistical test*

We will adopt here the most used framework of change in mean, and assume that the variance  $\sigma^2$  is known (or estimated).

Let  $q$  be the assumed number of changepoints in the dataset, we are interested in testing:

$$H_0 : q = 0$$

$$H_1 : q = m$$

To perform this test, we need a *test statistic* to be compared to a certain *threshold*. We will focus here on the likelihood ratio (LR) statistic evaluated at some time point  $\tau$ :

$$LR_\tau = \frac{1}{\sigma^2} \left[ \sum_{i=1}^n (X_i - \bar{X}_{1:n})^2 - \sum_{i=1}^n (X_i - \bar{X}_{1:\tau})^2 - \sum_{i=\tau+1}^n (X_i - \bar{X}_{\tau+1:n})^2 \right]$$

# Problem setting

## *Statistical test*

### CUSUM statistic

If we define  $C_\tau$  such that

$$C_\tau = \sqrt{\frac{\tau(n - \tau)}{n}} \bar{X}_{1:\tau} - \bar{X}_{\tau+1:n}$$

We note that  $LR_\tau = \frac{C_\tau^2}{\sigma^2}$



# Problem setting

## *Online/Offline inference*

### Online inference

The dataset keeps growing through time. We have a continuous flow of datapoints coming. The goal is still to estimate the localisation of the changepoint, but usually the quantity of data *after* the changepoint is limited. The data is analysed by flow.

#### Example

Seismographs continuously monitors the seismic activity in a given zone, stock market quotes are continuously updated.

### Offline inference

The size of the dataset is fixed, the aim is still to detect the changepoints, but this time we have access to possibly much more data after the changepoint.

#### Example

# What are we interested in ?

## *Single changepoint*

In some settings we know that there is at most one changepoint in the dataset. In this setting, the inference of the changepoint is well documented and an essentially complete theory exist.

We use the  $LR$  statistic and set a threshold  $c > 0$ . If the threshold is exceeded, than we conclude that there is a changepoint in the data and estimate this changepoint  $\tau$  with :

$$\hat{\tau} = \operatorname{argmax}_{\tau \in \{1, \dots, n-1\}} LR_{\tau} = \operatorname{argmax}_{\tau \in \{1, \dots, n-1\}} \frac{C_{\tau}^2}{\sigma^2}$$

### How to choose the threshold $c$ ?

Some theoretical results that ensure that the false positive rate tends to 0 as  $n \rightarrow \infty$  suggest that the threshold should be set to at least  $c = 2 \log n$ . However, this threshold can appear to be quite conservative in some applications, and some papers suggest to use  $c = 2 \log \log n$  instead.

# What are we interested in ?

## Multiple changepoints

- Unfortunately this framework becomes less efficient when there are several changepoints.
  - The presence of other changepoints affects the value of the CUSUM statistics, pushing it sometimes below the threshold even when there is an actual changepoint
  - This framework doesn't help to estimate the number of changepoints

### The penalised cost function

To deal with the possibility of multiple changepoints, we define the notion of cost function  $f(\tau_1, \dots, \tau_q | X_1, \dots, X_n)$ , which is a function that we aim to minimize in order to find the changepoints. One example of this function is simply the negative log-likelihood if we have a model for IID samples (e.g. Gaussian IID samples):

$$f(\tau_1, \dots, \tau_q | X_1, \dots, X_n) = \frac{n}{2} \log \sigma^2 + \frac{\sum_{j=1}^{q+1} \sum_{i=\tau_{j-1}+1}^{\tau_j} (X_i - \bar{X}_{\tau_{j-1}+1:\tau_j})^2}{2\sigma^2}$$

And aim to minimise the penalised cost:

$$\min_{\tau_1, \dots, \tau_q} f(\tau_1, \dots, \tau_q | X_1, \dots, X_n) + \beta(q)$$

# What are we interested in ?

## *Multiple changepoints*

### Example of penalisations

- BIC :  $\beta(q) = q \log(n)$
- AIC :  $\beta(q) = 2q$

Each type of penalisation has its own pros/cons and theoretical results. For instance, under some regularity condition, the BIC penalisation ensures that the inferred number of changepoints  $\hat{q}$  converges to the true number of changepoints  $q^*$ .

# What are we interested in ?

## *Multiple changepoints*

Another idea estimates multiple changepoints is to use a recursive approach. Since we have an efficient method to detect one changepoint, why not use this method iteratively on a succession of intervals ?

## Binary segmentation algorithm

### BS algorithm

- Find  $\hat{\tau}_1 = \underset{\tau \in \{1, \dots, n-1\}}{\operatorname{argmax}} C_\tau$ 
  - If  $C_{\hat{\tau}_1} > c_n$  then keep  $\hat{\tau}_1$  in memory and run the first step again on the intervals  $[1, \hat{\tau}_1]$  and  $[\hat{\tau}_1 + 1, n]$ .
  - If  $C_{\hat{\tau}_1} \leq c_n$  then stop the algorithm and return the list of the kept changepoints.

Or

- Find  $\hat{\tau}_1 = \underset{\tau \in \{1, \dots, n-1\}}{\operatorname{argmax}} LR_\tau$
- Run again the first step again on the intervals  $[1, \hat{\tau}_1]$  and  $[\hat{\tau}_1 + 1, n]$ , until the length of the interval is less than 2.
- Select the best changepoints collection using the penalised cost method.

# What are we interested in ?

## *Binary segmentation algorithm*


### Notes

A good algorithm

- Easy to implement (recursive algorithm)
- Computationally efficient (complexity in  $\mathcal{O}(n \log n)$ )

But

- Relies on the ability of the CUSUM statistic to detect one changepoint among possibly several changepoints...
- Need to select the right threshold  $c_n$  in the first version of the algorithm...



Multiple changepoints with the CUSUM statistic

# What are we interested in ?

## *Multiple changepoints*

### Wild binary segmentation (WBS) algorithm

To try to prevent the problem of ending up working in an interval that contains several changepoints and disturb the changepoint estimation, the Wild Binary Segmentation algorithm extend the BS algorithm to  $M$  random sub-intervals  $[s_m, e_m] \subset [1, n]$  and only keeping the highest CUSUM statistic among the batch of CUSUM statistic calculated. This way, we hope that at least one of the random sub-interval only contains one true changepoint.

#### BS algorithm

- Draw  $M$  sub-intervals of  $[1, n]$  and let  $\mathcal{M}$  be the set of the indices  $m$  such that  $[s_m, e_m] \subset [1, n]$ .
- Compute  $m_0, \hat{\tau}_1 = \underset{m \in \mathcal{M}, \tau \in [s_m, e_m]}{\operatorname{argmax}} C_\tau(X_{s_m:e_m})$ 
  - If  $LR_{\hat{\tau}_1} > c_n$  then keep  $\hat{\tau}_1$  in memory and run the first step again on the intervals  $[1, \hat{\tau}_1]$  and  $[\hat{\tau}_1 + 1, n]$ .
  - If  $LR_{\hat{\tau}_1} \leq c_n$  then stop the algorithm and return the list of the kept changepoints.

Or

- Draw  $M$  sub-intervals of  $[1, n]$  and let  $\mathcal{M}$  be the set of the indices  $m$  such that  $[s_m, e_m] \subset [1, n]$ .
- Compute  $m_0, \hat{\tau}_1 = \underset{m \in \mathcal{M}, \tau \in [s_m, e_m]}{\operatorname{argmax}} C_\tau(X_{s_m:e_m})$
- Run again the first step again on the intervals  $[1, \hat{\tau}_1]$  and  $[\hat{\tau}_1 + 1, n]$ , until the length of the interval is less then 2.
- Select the best changepoints collection using the penalised cost method.



# What are we interested in ?

## *Multiple changepoints*

### Wild binary segmentation (WBS) algorithm

#### Notes

A good algorithm

- Solve the initial problem of perturbation of the CUSUM statistics when there are several changepoints
- Computationally still quite efficient

But

- The computational cost increases with  $M$ .
- Need to select the right threshold  $c_n$  in the first version of the algorithm...

# What are we interested in ?

## *Multiple changepoints*

### Optimal partitionning

Let us come back to

$$\min_{\tau_1, \dots, \tau_q} f(\tau_1, \dots, \tau_q | X_1, \dots, X_n) + \beta(q)$$

This minimisation problem greatly simplifies if we assume that the criteria  $f$  can be broken down into a sum of costs over several segments:

$$f(\tau_1, \dots, \tau_q | X_1, \dots, X_n) = \sum_{j=1}^{q+1} c_{\tau_{j-1}, \tau_j}(X_{1:n})$$

Let us assume that  $\beta(q) = \lambda q$  where  $\lambda$  is a constant. We can therefore write the minimisation problem above as

$$Q_{n,\lambda}(q; \tau_1, \dots, \tau_q) = \sum_{j=1}^{q+1} c_{\tau_{j-1}, \tau_j}(X_{1:n}) + q\lambda$$

# What are we interested in ?

## *Multiple changepoints*

### Optimal partitionning

Now denote

$$Q_{t,\lambda} = \min_{q; \tau_1 < \dots < \tau_q < t} \sum_{j=1}^q c_{\tau_{j-1}, \tau_j}(X_{1:n}) + c_{\tau_q, t}(X_{1:n}) + q\lambda$$

This can be interpreted as the **minimum segmentation cost** of the time series between times  $0$  and  $t$ . It can be rewritten:

$$Q_{t,\lambda} = \min \left\{ c_{0,t}(X_{1:n}), \min_{\tau=1, \dots, t-1} Q_{\tau,\lambda} + c_{\tau,t}(X_{1:n}) + \lambda \right\}$$

Which simplifies further if we set  $Q_{0,\lambda} = -\lambda$  (this value is arbitrary as it represents the minimum segmenting cost of the series between  $0$  and  $0$ ):

# What are we interested in ?

## *Multiple changepoints*

### Optimal partitionning

This last expression has a recursive structure that one can exploit to solve this minimisation problem and estimate the changepoints  $(\hat{\tau}_j)_j$  and their number  $\hat{q}$ :

$$\hat{\tau}_1 = \underset{\tau=0,\dots,n-1}{\operatorname{argmin}} Q_{\tau,\lambda} + c_{\tau,n}(X_{1:n}) + \lambda$$

And then

$$\hat{\tau}_{j+1} = \underset{\tau=0,\dots,\hat{\tau}_j-1}{\operatorname{argmin}} Q_{\tau,\lambda} + c_{\tau,\hat{\tau}_j}(X_{1:n}) + \lambda$$

We recursively computes the  $(\hat{\tau}_j)_j$  until we find  $\hat{\tau} = 0$ .

# What are we interested in ?

## *Multiple changepoints*

### Pruned exact linear time (PELT) algorithm

The problem of the optimal partitioning is its computational cost:  $\mathcal{O}(n^2)$  (which can become prohibitive in some applications). The PELT algorithm has been introduced to reduce the computational cost of optimal partitioning by adding a *pruning* rule that prunes parts of the search space that are deemed to have a higher segmentation cost.

Indeed if at some point  $\tau$  we have:

$$Q_{\tau,\lambda} + c_{\tau+1:t}(X_{1:n}) + a > Q_t, \quad \text{for some } a > 0$$

Then  $\tau$  will never be an acceptable changepoint, and can thus be eliminated from the search space.

# Bayesian approach

## *Bayesian online changepoint detection*

This method, based on Adams and MacKay (2007), adopts a Bayesian approach to compute the posterior probability distribution of the “run length”  $r_t$  (i.e. the time after the last change point):

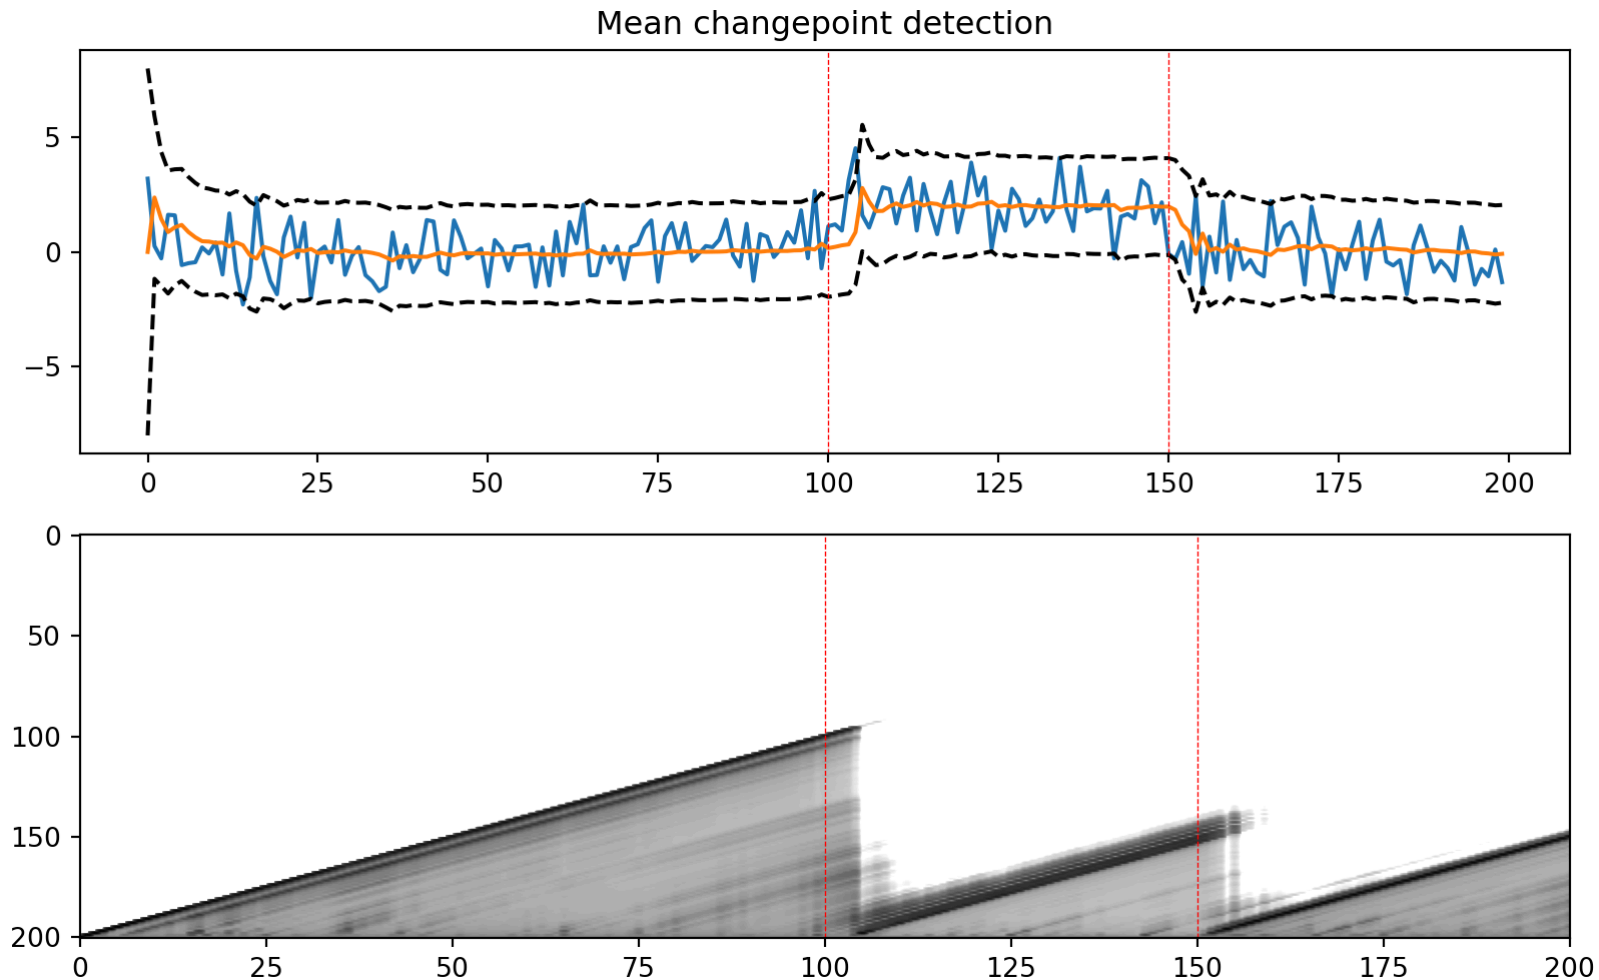
$$r_t = \begin{cases} 0, & \text{if } t \text{ is a changepoint} \\ r_{t-1} + 1, & \text{otherwise} \end{cases}$$

The idea is to assign a prior to the “hazard rate” (i.e. the frequency at which the changepoints occur) and use the exponential family posterior predictive closed formula to compute the posterior probability distribution of  $r_t$  at every  $t$ :

$$p(r_t | X_{1:t}) = \frac{p(r_t, X_{1:t})}{\sum_{r_{t'}} p(r_{t'}, X_{1:t})}$$

# Bayesian approach

## *Bayesian online changepoint detection*




Example of mean changepoint detection with the BOCD algorithm. The bottom plot represent the posterior probability distribution of  $r_t$ .



# Applications

# Canonical datasets


## *Well log*



Material response vs depth during an oil well drilling. This type of data is used to analyse the soil and determine what is the best depth and orientation for the forage.

# Canonical datasets

## *Dow Jones Industrial Average*




Daily returns of the Dow Jones Industrial Average index between 1970 and 2000. This can be one building block for market risk monitoring.

# Application in risk monitoring in Finance

## Dataset

To monitor financial markets and potentially take action to protect financial stability, trading venues monitor the volatility of the different stock index, to do so, they need some sort of *online* changepoint detection.



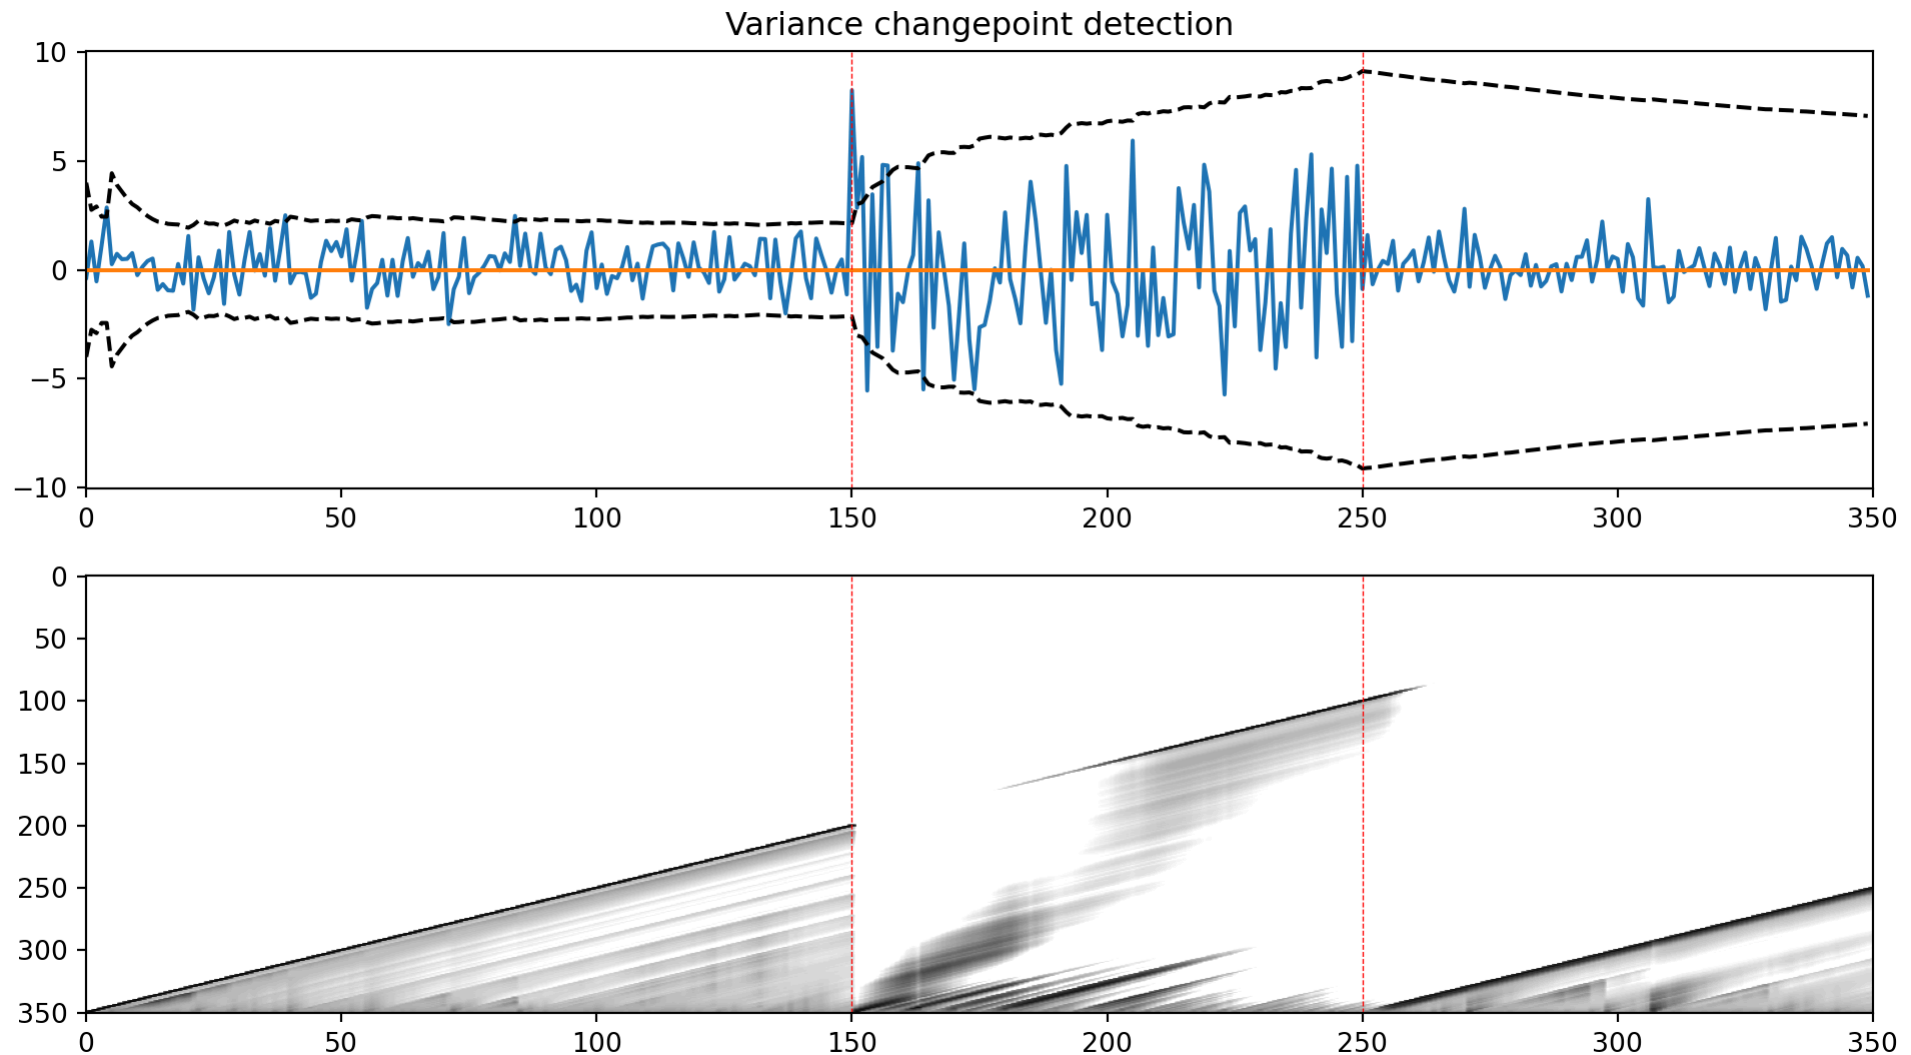
# Application in risk monitoring in Finance

## *Changepoint algorithm*

```
1 def BOCD_var(signal, alpha0, beta0, hazard_prob, mu = 0):
2
3     # Initialisation
4     T = len(signal)
5     log_R = -np.inf * np.ones((T+1, T+1)) # Run length posterior log-probability matrix
6     log_R[0, 0] = 0 # At time 0, the posterior probability is initialised to 1 at R = 0
7
8     posterior_precision = np.nan * np.empty(T) # Mean of the posterior distribution of the precision
9
10
11     log_message = np.array([0]) # message initialised at 1
12     log_H = np.log(hazard_prob) # Constant prior on changepoint probability.
13     log_1_minus_H = np.log(1-hazard_prob)
14
15     # Prior's parameters for the previous data point
16     prior_shape = np.array([alpha0])
17     prior_rate = np.array([beta0])
18
19
20     # Online posterior distribution of the run length update:
21     for t in range(1, T+1):
```

# Application in risk monitoring in Finance

## Results



Result of the BOCD algorithm on the variance changepoint dataset.

# Application in risk monitoring in Finance

## Results

### Remarks


Note that the inferred variance seem to do great at the beginning, but then lacks reactivity. This is beacause of the online nature of the BOCD algorithm: at every time-step it keeps memory of what happened in the past, and the longer the past is, the stronger the memory gets.

To prevent that, the *restarted* BOCD algorithm was developed by Alami, Maillard, and Feraud ([2020](#)) which restarts all over again every time a changepoint is detected.

# Other possible examples of applications

## *Retrospective analysis of GDP growth*

To analyse the effects of economic policies or the economic cycles, one can apply an *offline* changepoint analysis with a PELT algorithm to infer the number and the location of the changepoints, which can indicate a change in the economic cycle.





# Other possible examples of applications

## *Retrospective analysis of GDP growth*

```
1 import ruptures as rpt
2
3 data = pd.read_csv("../Datasets/gdp.csv")
4
5 country_gdp = data.loc[data["Country Code"] == "USA"].melt(id_vars="Country Name", value_vars = [str(i) for i in range(1950, 2015)])
6 country_gdp["Growth"] = country_gdp["GDP"].pct_change()
7
8 growth_data = np.array(country_gdp['Growth'])[1:]*100
9
10 algo = rpt.Pelt(model = "l1", min_size = 2).fit(growth_data)
11 changepoints = algo.predict(pen = np.log(len(growth_data)))
12
13
14 plt.figure(figsize = (12, 5))
15 plt.plot(np.array(country_gdp["Year"], dtype = float)[1:], growth_data, label = "USA")
16 plt.vlines([np.array(country_gdp["Year"], dtype = float)[1:][i] for i in changepoints[:-1]], linestyle = "dashed")
17 plt.xticks(rotation = 45);
18 plt.ylim((1.05*min(growth_data), 1.05*max(growth_data)));
19 plt.xlabel("Year")
20 plt.legend()
21 plt.ylabel("GDP growth (%)");
```

# Useful libraries, package and repository

## In R

- `mcp`
- `segmented`

## In Python

- `ruptures`
- `changeponyt`

## Useful GitHub repositories

<https://github.com/gwgundersen/bocd> <https://github.com/alan-turing-institute/TCPD/tree/master>

# References

## *Code and presentation*

GitHub repository

## *Textbooks and articles*

Adams, Ryan Prescott, and David J. C. MacKay. 2007. “Bayesian Online Changepoint Detection.” arXiv. <https://doi.org/10.48550/arXiv.0710.3742>.

Alami, Reda, Odalric Maillard, and Raphael Feraud. 2020. “Restarted Bayesian Online Change-Point Detector Achieves Optimal Detection Delay.” In *Proceedings of the 37th International Conference on Machine Learning*, edited by Hal Daume III and Aarti Singh, 119:211–21. Proceedings of Machine Learning Research. PMLR. <https://proceedings.mlr.press/v119/alami20a.html>.

“Bayesian Online Changepoint Detection.” n.d. Accessed June 6, 2024. <https://gregorygundersen.com/blog/2019/08/13/bocd/#adams2007bayesian>.

Fearnhead, Paul, and Piotr Fryzlewicz. 2022. “Detecting A Single Change-Point.” arXiv. <http://arxiv.org/abs/2210.07066>.

———. 2024. “The Multiple Change-in-Gaussian-Mean Problem.” arXiv. <http://arxiv.org/abs/2405.06796>.

Feasel, Kevin. 2022. *Finding Ghosts in Your Data: Anomaly Detection Techniques with Examples in Python*. Berkeley, CA: Apress. <https://doi.org/10.1007/978-1-4842-8870-2>.

